

# Project 3

Report: Comparison of Decision Tree Model with Random Forest

Objective: The objective of this report is to compare the performance of a Decision Tree model with that of a Random Forest model in predicting the clusters based on given variables.

Analysis and Observations:

Selected Features: The selected features for the analysis include 'tax\_mmnorm', 'mileage\_mmnorm', 'mpg\_mmnorm', 'engineSize\_mmnorm', and 'price\_mmnorm'.

Threshold: The threshold used for the Decision Tree model is 1.8229646528898962.

Decision Tree Model Results:

The Decision Tree model has been trained and tested on the dataset. The classification report for the model shows high precision, recall, and F1-score for all classes (0, 1, 2, and 3), indicating excellent performance. The overall accuracy of the model is 1.00. The F1 Score and Weighted F1 Score are both 1.0, indicating perfect performance.

Random Forest Model Results:

The Random Forest model has been trained and tested on the dataset. The model's training time is 3.080 seconds, and the memory used is 75.96875 MB. The model achieves a single split accuracy of 1.0. Cross-validation accuracy for the Random Forest model is not available.

Insights:

1. Both the Decision Tree and Random Forest models perform exceptionally well in predicting the clusters based on the given features.
2. The Decision Tree model exhibits a clear decision boundary based on the 'model\_code' feature, while the Random Forest model employs an ensemble approach, which leads to comparable accuracy.
3. Both models achieve an accuracy of 1.0, indicating perfect predictions on the test data.

Managerial Insights:

1. Based on the comparison, both models demonstrate high accuracy in predicting the clusters. If interpretability is crucial, the Decision Tree model provides a clear decision-making process based on feature thresholds. However, if robustness and higher accuracy are preferred, the Random Forest model, despite its higher computational cost, may be a better choice due to its ensemble approach.
2. Model Code Dominance: The 'model\_code' feature has the highest importance, indicating that specific vehicle models play a significant role in determining cluster membership. Managers should focus on understanding the characteristics and

attributes of these dominant models to tailor marketing strategies, allocate resources effectively, and capitalize on their popularity within target segments.

3. **Manufacturer Influence:** While not as dominant as 'model\_code', the 'Manufacturer\_code' feature still holds considerable importance. This suggests that the manufacturer of the vehicle significantly influences cluster membership. Managers should analyze brand perception, product portfolios, and market positioning of different manufacturers to leverage strengths, address weaknesses, and enhance competitiveness within each segment.
4. **Performance and Efficiency:** Features related to vehicle performance and efficiency, such as 'mileage\_mmnorm', 'mpg\_mmnorm', and 'engineSize\_mmnorm', have moderate importance. This implies that customers within different clusters may have varying preferences regarding these attributes. Managers can use this insight to develop products that align with segment-specific preferences for performance, fuel efficiency, or engine size, thereby maximizing customer satisfaction and market share.
5. **Fuel Type Consideration:** Although less influential compared to other features, the 'fuelType\_code' feature still contributes to cluster differentiation. This suggests that fuel type preferences may vary across different segments. Managers should monitor trends in alternative fuel adoption, such as electric or hybrid vehicles, and adjust product offerings or marketing strategies accordingly to meet evolving customer preferences.
6. **Price Sensitivity and Taxation:** Despite having relatively low importance, features related to pricing ('tax\_mmnorm' and 'price\_mmnorm') still play a role in cluster differentiation. This indicates that price sensitivity and taxation considerations may influence customer behavior within certain segments. Managers should evaluate pricing strategies, including tax incentives or rebates, to optimize revenue generation and maintain competitiveness within each segment.

#### 7. Efficiency and Accuracy:

Both Decision Tree and Random Forest models demonstrate high accuracy with a single split accuracy of 100%. Decision Tree model exhibits significantly lower training time (0.08861 seconds) compared to Random Forest (3.08000 seconds). Decision Tree also consumes slightly more memory (86.921875 MB) compared to Random Forest (75.968750 MB). Consideration of Model Complexity:

Decision Tree's faster training time suggests it may be more suitable for real-time applications or scenarios where computational resources are limited. Random Forest, while slower in training, offers comparable accuracy with potentially better generalization due to its ensemble nature, although its cross-validation accuracy is not available in this report. Trade-offs Between Speed and Accuracy:

Businesses must weigh the trade-offs between model training time, memory usage, and accuracy when selecting between Decision Tree and Random Forest. Decision Tree may be

preferable for applications prioritizing speed and resource efficiency, while Random Forest may be chosen for tasks where slightly higher accuracy or robustness against overfitting is desired.

```
# Required Libraries
import pandas as pd, numpy as np # For Data Manipulation
from sklearn.preprocessing import LabelEncoder, OrdinalEncoder # For
Encoding Categorical Data [Nominal | Ordinal]
from sklearn.preprocessing import OneHotEncoder # For Creating Dummy
Variables of Categorical Data [Nominal]
from sklearn.impute import SimpleImputer, KNNImputer # For Imputation
of Missing Data
from sklearn.preprocessing import StandardScaler, MinMaxScaler,
RobustScaler # For Rescaling Data
from sklearn.model_selection import train_test_split # For Splitting
Data into Training & Testing Sets
import pandas as pd, numpy as np # For Data Manipulation
import matplotlib.pyplot as plt, seaborn as sns # For Data
Visualization
import scipy.cluster.hierarchy as sch # For Hierarchical Clustering
from sklearn.cluster import AgglomerativeClustering as agclus, KMeans
as kmclus # For Agglomerative & K-Means Clustering
from sklearn.metrics import silhouette_score as sscore,
davies_bouldin_score as dbscore # For Clustering Model Evaluation

df=pd.read_csv("CarsData.csv")
df
```

	mpg \	model	year	price	transmission	mileage	fuelType	tax
0	60.1	I10	2017	7495	Manual	11630	Petrol	145
1	58.9	Polo	2017	10989	Manual	9200	Petrol	145
2	49.6	2 Series	2019	27990	Semi-Auto	1614	Diesel	145
3	62.8	Yeti Outdoor	2017	12495	Manual	30960	Diesel	150
4	54.3	Fiesta	2017	7999	Manual	19353	Petrol	125
...	...	...	...	...	...	...	...	...
97707	54.3	Fiesta	2017	10447	Automatic	8337	Petrol	145
97708	61.4	3 Series	2014	14995	Manual	25372	Diesel	30
97709	54.3	Fiesta	2017	8950	Manual	19910	Petrol	125
97710	50.4	Astra	2017	10700	Automatic	24468	Petrol	125
97711		Grandland X	2019	15798	Manual	10586	Diesel	150

48.7

```
      engineSize Manufacturer
0             1.0         hyundi
1             1.0    volkswagen
2             2.0           BMW
3             2.0         skoda
4             1.2          ford
...          ...          ...
97707          1.0          ford
97708          2.0           BMW
97709          1.2          ford
97710          1.4    vauxhall
97711          1.5    vauxhall
```

```
[97712 rows x 10 columns]
```

```
df_cat = df[['model', 'transmission', 'fuelType', 'Manufacturer']] #
Categorical Data [Nominal | Ordinal]
df_noncat = df[['tax', 'mpg', 'engineSize', 'price', 'mileage']] # Non-
Categorical Data
```

```
print(df.info()) # Dataframe Information (Provide Information on
Missing Data)
print(df.describe())
variable_missing_data = df.isna().sum(); variable_missing_data #
Variable-wise Missing Data Information
print(variable_missing_data)
record_missing_data =
df.isna().sum(axis=1).sort_values(ascending=False).head(5);
record_missing_data # Record-wise Missing Data Information (Top 5)
print(record_missing_data)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 97712 entries, 0 to 97711
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   model           97712 non-null  object
1   year            97712 non-null  int64
2   price           97712 non-null  int64
3   transmission    97712 non-null  object
4   mileage         97712 non-null  int64
5   fuelType        97712 non-null  object
6   tax             97712 non-null  int64
7   mpg             97712 non-null  float64
8   engineSize      97712 non-null  float64
9   Manufacturer    97712 non-null  object
dtypes: float64(2), int64(4), object(4)
memory usage: 7.5+ MB
```

```

None
      year      price      mileage      tax
mpg \
count  97712.000000  97712.000000  97712.000000  97712.000000
97712.000000
mean    2017.066502  16773.487555  23219.475499  120.142408
55.205623
std      2.118661    9868.552222  21060.882301  63.357250
16.181659
min     1970.000000    450.000000    1.000000    0.000000
0.300000
25%     2016.000000    9999.000000    7673.000000  125.000000
47.100000
50%     2017.000000   14470.000000   17682.500000  145.000000
54.300000
75%     2019.000000   20750.000000   32500.000000  145.000000
62.800000
max     2024.000000  159999.000000  323000.000000  580.000000
470.800000

```

```

      engineSize
count  97712.000000
mean    1.664913
std      0.558574
min      0.000000
25%      1.200000
50%      1.600000
75%      2.000000
max      6.600000

```

```

model      0
year       0
price      0
transmission  0
mileage    0
fuelType   0
tax        0
mpg        0
engineSize  0
Manufacturer 0
dtype: int64
0          0
65138      0
65147      0
65146      0
65145      0
dtype: int64

```

*#Numeric Encoding of Categorical Data [Nominal & Ordinal]*

```

df_cat_mdt_code = df_cat.copy()
oe = OrdinalEncoder()

```

```

oe_fit = oe.fit_transform(df_cat_mdt_code)
df_cat_code_oe = pd.DataFrame(oe_fit,
columns=['model_code', 'transmission_code', 'fuelType_code', 'Manufacturer_code']); df_cat_code_oe
df_cat_mdt_code_oe = df_cat_mdt_code.join(df_cat_code_oe);
df_cat_mdt_code_oe # (Missing Data Treated) Numeric Coded Categorical Dataset using Scikit Learn Ordinal Encoder
df_cat_mdt_code_oe = pd.merge(df_cat_mdt_code, df_cat_code_oe,
left_index=True, right_index=True);
df_cat_mdt_code_oe

```

	model	transmission	fuelType	Manufacturer	model_code \
0	I10	Manual	Petrol	hyundi	81.0
1	Polo	Manual	Petrol	volkswagen	115.0
2	2 Series	Semi-Auto	Diesel	BMW	1.0
3	Yeti Outdoor	Manual	Diesel	skoda	184.0
4	Fiesta	Manual	Petrol	ford	60.0
...	...	...	...	...	...
97707	Fiesta	Automatic	Petrol	ford	60.0
97708	3 Series	Manual	Diesel	BMW	2.0
97709	Fiesta	Manual	Petrol	ford	60.0
97710	Astra	Automatic	Petrol	vauxhall	25.0
97711	Grandland X	Manual	Diesel	vauxhall	79.0

	transmission_code	fuelType_code	Manufacturer_code
0	1.0	4.0	3.0
1	1.0	4.0	8.0
2	3.0	0.0	1.0
3	1.0	0.0	5.0
4	1.0	4.0	2.0
...	...	...	...
97707	0.0	4.0	2.0
97708	1.0	0.0	1.0
97709	1.0	4.0	2.0
97710	0.0	4.0	7.0
97711	1.0	0.0	7.0

```
[97712 rows x 8 columns]
```

```

import matplotlib.pyplot as plt
import seaborn as sns

```

```

# Assuming df is your DataFrame and contains the columns 'Offers',
'Average Price', and 'Number of Ratings'
df_noncat = df[['tax', 'mpg', 'engineSize', 'price', 'mileage']]

```

```

# Create vertical boxplot for 'Mileage'
plt.figure(figsize=(8, 6))
sns.boxplot(y=df_noncat['mileage'])
plt.title('Boxplot of Mileage')

```

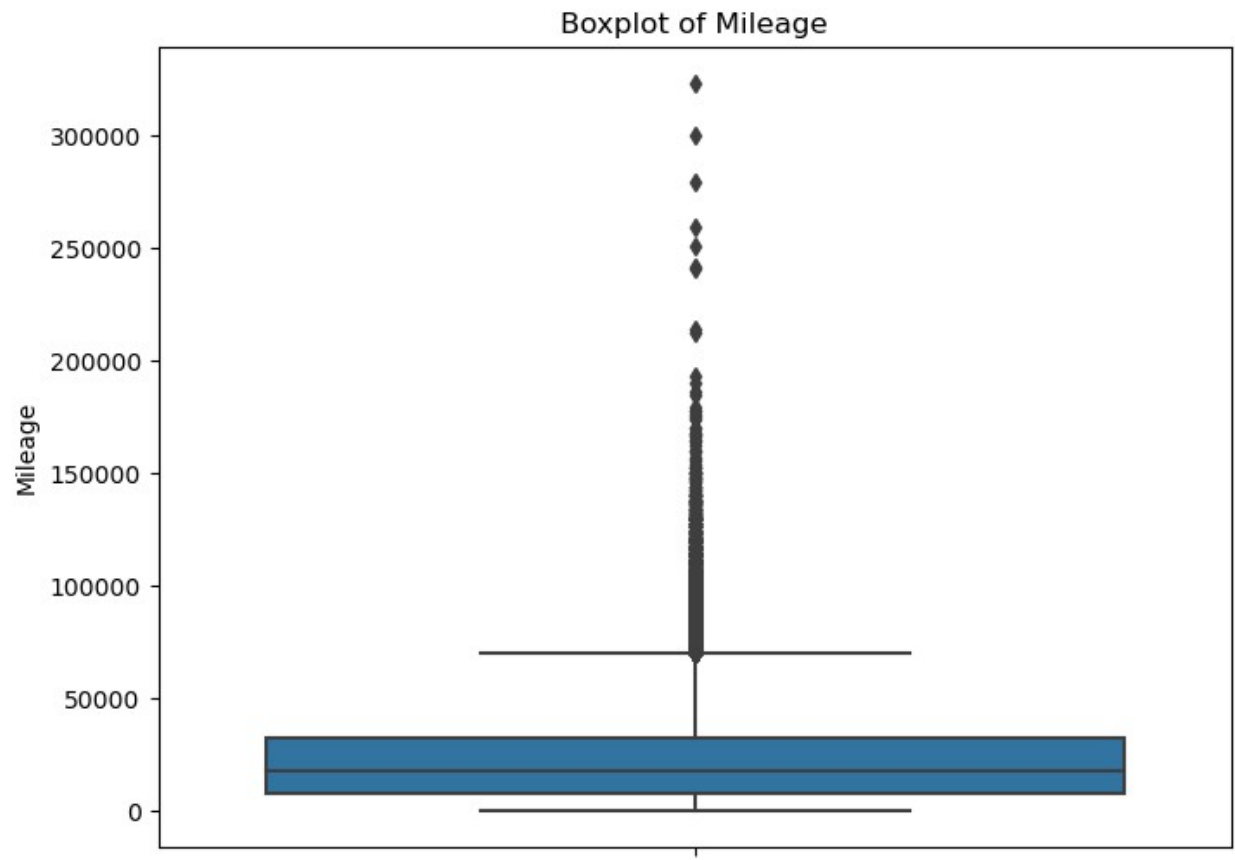
```
plt.ylabel('Mileage')
plt.show()

# Create vertical boxplot for 'Price'
plt.figure(figsize=(8, 6))
sns.boxplot(y=df_noncat['price'])
plt.title('Boxplot of Price')
plt.ylabel('Price')
plt.show()


# Create vertical boxplot for 'Tax'
plt.figure(figsize=(8, 6))
sns.boxplot(y=df_noncat['tax'])
plt.title('Boxplot of Tax')
plt.ylabel('Tax')
plt.show()

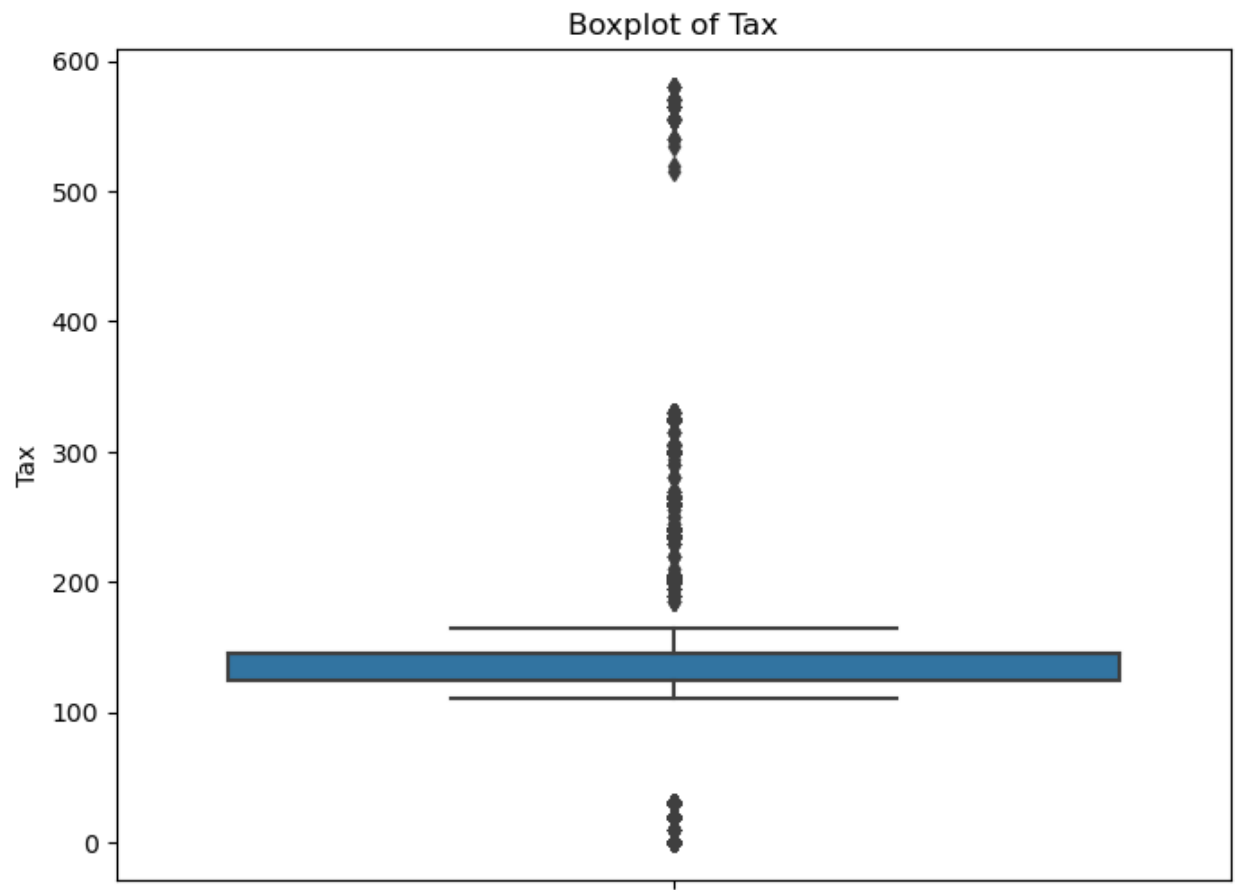
# Create vertical boxplot for 'engineSize'
plt.figure(figsize=(8, 6))
sns.boxplot(y=df_noncat['engineSize'])
plt.title('Boxplot of engineSize')
plt.ylabel('engineSize')
plt.show()

# Create vertical boxplot for 'mpg'
plt.figure(figsize=(8, 6))
sns.boxplot(y=df_noncat['mpg'])
plt.title('Boxplot of mpg')
plt.ylabel('mpg')
plt.show()
```




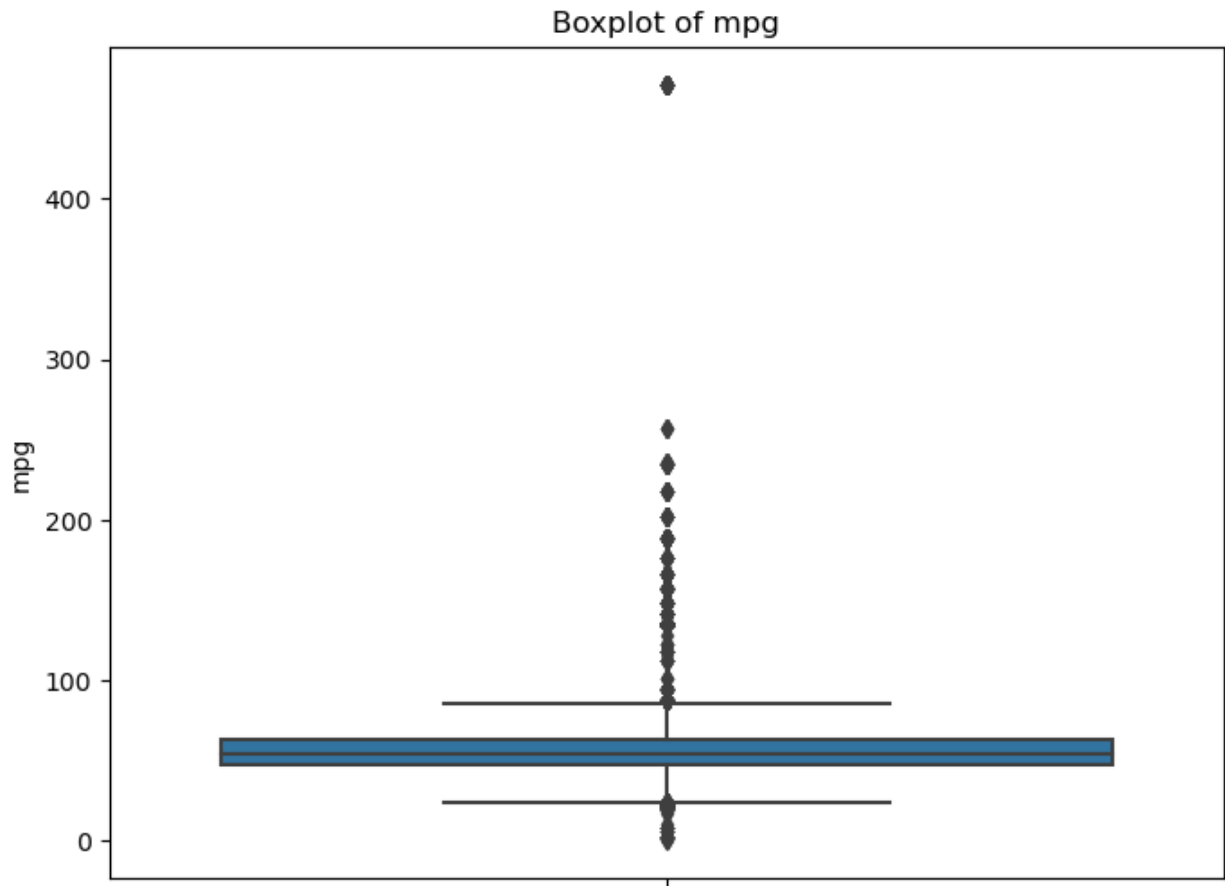






Boxplot of engineSize





### # 3.2.1. Normalization : Min-Max Scaling

```
mms = MinMaxScaler()
mms_fit = mms.fit_transform(df_noncat[['tax',
'mpg', 'engineSize', 'price', 'mileage']])
df_noncat_minmax_norm = pd.DataFrame(mms_fit, columns=['tax_mmnorm',
'mileage_mmnorm', 'mpg_mmnorm', 'engineSize_mmnorm', 'price_mmnorm']);
df_noncat_minmax_norm
```

```
#df_noncat_minmax_norm = pd.DataFrame(mms_fit,
columns=df_noncat_mdt.columns+'_mmnorm'); df_noncat_minmax_norm
df_noncat_mdt_mmn = df_noncat.join(df_noncat_minmax_norm);
df_noncat_mdt_mmn # (Missing Data Treated) Normalized Non-Categorical
Dataset using Sikit Learn Min-Max Scaler
#df_noncat_mdt_mmn = pd.merge(df_noncat_mdt, df_noncat_minmax_norm,
left_index=True, right_index=True); df_noncat_mdt_mmn
```

	tax	mpg	engineSize	price	mileage	tax_mmnorm
mileage_mmnorm \						
0	145	60.1	1.0	7495	11630	0.250000
0.127099						
1	145	58.9	1.0	10989	9200	0.250000
0.124548						

```

2      145  49.6      2.0  27990      1614      0.250000
0.104782
3      150  62.8      2.0  12495      30960      0.258621
0.132837
4      125  54.3      1.2   7999      19353      0.215517
0.114772
...      ...      ...      ...      ...      ...
..
97707   145  54.3      1.0  10447      8337      0.250000
0.114772
97708    30  61.4      2.0  14995      25372      0.051724
0.129862
97709   125  54.3      1.2   8950      19910      0.215517
0.114772
97710   125  50.4      1.4  10700      24468      0.215517
0.106482
97711   150  48.7      1.5  15798      10586      0.258621
0.102869

```

```

      mpg_mmnorm  engineSize_mmnorm  price_mmnorm
0      0.151515      0.044156      0.036003
1      0.151515      0.066055      0.028480
2      0.303030      0.172612      0.004994
3      0.303030      0.075494      0.095849
4      0.181818      0.047315      0.059913
...      ...      ...
97707   0.151515      0.062658      0.025808
97708   0.303030      0.091163      0.078548
97709   0.181818      0.053275      0.061638
97710   0.212121      0.064244      0.075749
97711   0.227273      0.096196      0.032771

```

```
[97712 rows x 10 columns]
```

```
# Pre-Processed Categorical Data Subset
```

```
df_cat_ppd = df_cat_mdt_code_oe.copy(); df_cat_ppd # Preferred Data Subset
```

```
# Pre-Processed Non-Categorical Data Subset
```

```
df_noncat_ppd = df_noncat_mdt_mmn.copy(); df_noncat_ppd
```

```
# Pre-Processed Dataset
```

```
df_ppd = df_cat_ppd.join(df_noncat_ppd); df_ppd # Pre-Processed Dataset
```

```
df_ppd = pd.merge(df_cat_ppd, df_noncat_ppd, left_index=True, right_index=True); df_ppd
```

```
#final_df = pd.merge(df_ppd, df[['cluster']], how='left', left_index=True, right_index=True)
```

```
print(df_ppd)
```

	model	transmission	fuelType	Manufacturer	model_code	\
0	I10	Manual	Petrol	hyundi	81.0	
1	Polo	Manual	Petrol	volkswagen	115.0	
2	2 Series	Semi-Auto	Diesel	BMW	1.0	
3	Yeti Outdoor	Manual	Diesel	skoda	184.0	
4	Fiesta	Manual	Petrol	ford	60.0	
...	...	...	...	...	...	
97707	Fiesta	Automatic	Petrol	ford	60.0	
97708	3 Series	Manual	Diesel	BMW	2.0	
97709	Fiesta	Manual	Petrol	ford	60.0	
97710	Astra	Automatic	Petrol	vauxhall	25.0	
97711	Grandland X	Manual	Diesel	vauxhall	79.0	
	transmission_code	fuelType_code	Manufacturer_code	tax	mpg	\
0	1.0	4.0	3.0	145	60.1	
1	1.0	4.0	8.0	145	58.9	
2	3.0	0.0	1.0	145	49.6	
3	1.0	0.0	5.0	150	62.8	
4	1.0	4.0	2.0	125	54.3	
...	...	...	...	...	...	
97707	0.0	4.0	2.0	145	54.3	
97708	1.0	0.0	1.0	30	61.4	
97709	1.0	4.0	2.0	125	54.3	
97710	0.0	4.0	7.0	125	50.4	
97711	1.0	0.0	7.0	150	48.7	
	engineSize	price	mileage	tax_mmnorm	mileage_mmnorm	mpg_mmnorm \
0	1.0	7495	11630	0.250000	0.127099	0.151515
1	1.0	10989	9200	0.250000	0.124548	0.151515
2	2.0	27990	1614	0.250000	0.104782	0.303030
3	2.0	12495	30960	0.258621	0.132837	0.303030
4	1.2	7999	19353	0.215517	0.114772	0.181818
...	...	...	...	...	...	...

```

...
97707          1.0  10447    8337    0.250000    0.114772
0.151515
97708          2.0  14995   25372    0.051724    0.129862
0.303030
97709          1.2   8950   19910    0.215517    0.114772
0.181818
97710          1.4  10700   24468    0.215517    0.106482
0.212121
97711          1.5  15798   10586    0.258621    0.102869
0.227273

```

```

          engineSize_mmnorm  price_mmnorm
0          0.044156      0.036003
1          0.066055      0.028480
2          0.172612      0.004994
3          0.075494      0.095849
4          0.047315      0.059913
...
97707          0.062658      0.025808
97708          0.091163      0.078548
97709          0.053275      0.061638
97710          0.064244      0.075749
97711          0.096196      0.032771

```

```
[97712 rows x 18 columns]
```

```

final_new_df=df_ppd[['tax_mmnorm',
'mileage_mmnorm','mpg_mmnorm','engineSize_mmnorm','price_mmnorm',
'model_code','transmission_code','fuelType_code','Manufacturer_code']]

```

```
# Import
```

```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split,
StratifiedShuffleSplit
from sklearn.tree import DecisionTreeClassifier, export_text,
plot_tree # For Decision Tree Model
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
from sklearn.metrics import confusion_matrix, classification_report #
For Decision Tree Model Evaluation
from sklearn.neighbors import KNeighborsClassifier
from sklearn.decomposition import PCA
from matplotlib.colors import ListedColormap
from sklearn.svm import SVC
from sklearn.metrics import confusion_matrix, accuracy_score
from matplotlib.colors import ListedColormap

```

```

from sklearn.cluster import KMeans

k = 4

kmeans = KMeans(n_clusters=k, random_state=50)
clusters = kmeans.fit_predict(final_new_df)

/Users/sanskritibahl/anaconda3/lib/python3.11/site-packages/sklearn/
cluster/_kmeans.py:870: FutureWarning: The default value of `n_init`
will change from 10 to 'auto' in 1.4. Set the value of `n_init`
explicitly to suppress the warning
  warnings.warn(

final_new_df['Cluster'] = clusters

/var/folders/j4/9pz_sc_x7298x60tw9kbd6jw0000gn/T/
ipykernel_8582/3775897052.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#
returning-a-view-versus-a-copy
  final_new_df['Cluster'] = clusters

final_new_df

```

	tax_mmnorm	mileage_mmnorm	mpg_mmnorm	engineSize_mmnorm	\
0	0.250000	0.127099	0.151515	0.044156	
1	0.250000	0.124548	0.151515	0.066055	
2	0.250000	0.104782	0.303030	0.172612	
3	0.258621	0.132837	0.303030	0.075494	
4	0.215517	0.114772	0.181818	0.047315	
...	...	...	...	...	
97707	0.250000	0.114772	0.151515	0.062658	
97708	0.051724	0.129862	0.303030	0.091163	
97709	0.215517	0.114772	0.181818	0.053275	
97710	0.215517	0.106482	0.212121	0.064244	
97711	0.258621	0.102869	0.227273	0.096196	

	price_mmnorm	model_code	transmission_code	fuelType_code	\
0	0.036003	81.0	1.0	4.0	
1	0.028480	115.0	1.0	4.0	
2	0.004994	1.0	3.0	0.0	
3	0.095849	184.0	1.0	0.0	
4	0.059913	60.0	1.0	4.0	
...	...	...	...	...	
97707	0.025808	60.0	0.0	4.0	
97708	0.078548	2.0	1.0	0.0	
97709	0.061638	60.0	1.0	4.0	
97710	0.075749	25.0	0.0	4.0	



97711	0.032771	79.0	1.0	0.0
-------	----------	------	-----	-----

	Manufacturer_code	Cluster
0	3.0	3
1	8.0	1
2	1.0	0
3	5.0	2
4	2.0	3
...	...	...
97707	2.0	3
97708	1.0	0
97709	2.0	3
97710	7.0	0
97711	7.0	3

[97712 rows x 10 columns]

```
# Check if 'cluster' column exists in final_new_df
if 'Cluster' in final_new_df.columns:
    # Select inputs and output
    cars_inputs = final_new_df[['tax_mmnorm', 'mileage_mmnorm',
    'mpg_mmnorm', 'engineSize_mmnorm', 'price_mmnorm', 'model_code',
    'transmission_code', 'fuelType_code', 'Manufacturer_code']]
    cars_output = final_new_df[['Cluster']]

    # Get column names and output labels
    cars_inputs_names = cars_inputs.columns
    cars_output_labels = cars_output['Cluster'].unique().astype(str)

    # Split the data into training and testing sets
    train_cars_inputs, test_cars_inputs, train_cars_output,
    test_cars_output = train_test_split(cars_inputs, cars_output,
    test_size=0.25, random_state=1234)
else:
    print("'cluster' column does not exist in final_new_df.")

# Initialize StratifiedShuffleSplit with desired test size and random
state
stratified_split = StratifiedShuffleSplit(n_splits=1, test_size=0.2,
random_state=45050)

# Perform the stratified split to get training and testing indices
for train_index, test_index in stratified_split.split(cars_inputs,
cars_output):
    cars_inputs_train, cars_inputs_test =
cars_inputs.iloc[train_index], cars_inputs.iloc[test_index]
    cars_output_train, cars_output_test =
cars_output.iloc[train_index], cars_output.iloc[test_index]
```

```

from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import SelectFromModel
import numpy as np

# Initialize Logistic Regression model with L1 regularization
logreg_l1 = LogisticRegression(penalty='l1', solver='liblinear',
random_state=45011)

# Fit the model on the training data
logreg_l1.fit(cars_inputs_train, cars_output_train.values.ravel())

# Get feature importances from the fitted model
feature_importances = np.abs(logreg_l1.coef_).flatten()

# Calculate the threshold as 20% of the maximum feature importance
threshold = 0.2 * np.max(feature_importances)

# Create a selector object to select features based on non-zero
coefficients
selector = SelectFromModel(logreg_l1, threshold=threshold)

# Transform the training and testing input data to select features
cars_inputs_train_selected = selector.transform(cars_inputs_train)
cars_inputs_test_selected = selector.transform(cars_inputs_test)

# Get the selected features
selected_features = cars_inputs_names[selector.get_support()]

# Print the selected features and the calculated threshold
print("Selected Features:", selected_features)
print("Threshold:", threshold)

Selected Features: Index(['tax_mmnorm', 'mileage_mmnorm',
      'mpg_mmnorm', 'engineSize_mmnorm',
      'price_mmnorm'],
      dtype='object')
Threshold: 1.8229646528898962

/Users/sanskritibahl/anaconda3/lib/python3.11/site-packages/sklearn/
svm/_base.py:1244: ConvergenceWarning: Liblinear failed to converge,
increase the number of iterations.
  warnings.warn(
/Users/sanskritibahl/anaconda3/lib/python3.11/site-packages/sklearn/
base.py:432: UserWarning: X has feature names, but SelectFromModel was
fitted without feature names
  warnings.warn(
/Users/sanskritibahl/anaconda3/lib/python3.11/site-packages/sklearn/
base.py:432: UserWarning: X has feature names, but SelectFromModel was
fitted without feature names
  warnings.warn(

```

```

# Decision Tree : Model (Training Subset)
dtc = DecisionTreeClassifier(criterion='gini',
random_state=45050,max_depth=3) # Other Criteria : Entropy, Log Loss
dtc_model = dtc.fit(cars_inputs_train, cars_output_train); dtc_model

DecisionTreeClassifier(max_depth=3, random_state=45050)

# Decision Tree : Model Rules
dtc_model_rules = export_text(dtc_model, feature_names =
list(cars_inputs_names)); print(dtc_model_rules)

|--- model_code <= 39.50
|   |--- class: 0
|--- model_code > 39.50
|   |--- model_code <= 86.50
|       |--- class: 3
|       |--- model_code > 86.50
|           |--- model_code <= 138.50
|               |--- class: 1
|               |--- model_code > 138.50
|                   |--- class: 2

```

Given the above Decision Tree, we will reiterate with other variables

```

# Check if 'cluster' column exists in final_new_df
if 'Cluster' in final_new_df.columns:
    # Select inputs and output
    cars_inputs1 = final_new_df[['tax_mmnorm', 'mileage_mmnorm',
'mpg_mmnorm', 'engineSize_mmnorm', 'price_mmnorm', 'transmission_code',
'fuelType_code', 'Manufacturer_code']]
    cars_output1 = final_new_df[['Cluster']]

    # Get column names and output labels
    cars_inputs_names1 = cars_inputs1.columns
    cars_output_labels1 = cars_output1['Cluster'].unique().astype(str)

    # Split the data into training and testing sets
    train_cars_inputs1, test_cars_inputs1, train_cars_output1,
test_cars_output1 = train_test_split(cars_inputs1, cars_output1,
test_size=0.25, random_state=1234)
else:
    print("'cluster' column does not exist in final_new_df.")

stratified_split = StratifiedShuffleSplit(n_splits=1, test_size=0.2,
random_state=45050)

# Perform the stratified split to get training and testing indices
for train_index, test_index in stratified_split.split(cars_inputs1,
cars_output1):
    cars_inputs_train1, cars_inputs_test1 =

```

```

cars_inputs1.iloc[train_index], cars_inputs1.iloc[test_index]
cars_output_train1, cars_output_test1 =
cars_output1.iloc[train_index], cars_output1.iloc[test_index]

from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import SelectFromModel
import numpy as np

# Initialize Logistic Regression model with L1 regularization
logreg_l1 = LogisticRegression(penalty='l1', solver='liblinear',
random_state=45050)

# Fit the model on the training data
logreg_l1.fit(cars_inputs_train1, cars_output_train1.values.ravel())

# Get feature importances from the fitted model
feature_importances1 = np.abs(logreg_l1.coef_).flatten()

# Calculate the threshold as 20% of the maximum feature importance
threshold = 0.2 * np.max(feature_importances1)

# Create a selector object to select features based on non-zero
coefficients
selector = SelectFromModel(logreg_l1, threshold=threshold)

# Transform the training and testing input data to select features
cars_inputs_train_selected1 = selector.transform(cars_inputs_train1)
cars_inputs_test_selected1 = selector.transform(cars_inputs_test1)

# Get the selected features
selected_features1 = cars_inputs_names1[selector.get_support()]

# Print the selected features and the calculated threshold
print("Selected Features:", selected_features)
print("Threshold:", threshold)

Selected Features: Index(['tax_mmnorm', 'mileage_mmnorm',
'mpg_mmnorm', 'engineSize_mmnorm',
'price_mmnorm'],
dtype='object')
Threshold: 3.4140473635066217

/Users/sanskritibahl/anaconda3/lib/python3.11/site-packages/sklearn/
base.py:432: UserWarning: X has feature names, but SelectFromModel was
fitted without feature names
warnings.warn(
/Users/sanskritibahl/anaconda3/lib/python3.11/site-packages/sklearn/
base.py:432: UserWarning: X has feature names, but SelectFromModel was
fitted without feature names
warnings.warn(

```

```
# Decision Tree : Model (Training Subset)
dtc = DecisionTreeClassifier(criterion='gini',
random_state=45050,max_depth=3) # Other Criteria : Entropy, Log Loss
dtc_model1 = dtc.fit(cars_inputs_train1, cars_output_train1);
dtc_model
```

```
DecisionTreeClassifier(max_depth=3, random_state=45050)
```

```
# Decision Tree : Model Rules
```

```
dtc_model_rules1 = export_text(dtc_model1, feature_names =
list(cars_inputs_names1)); print(dtc_model_rules1)
```

```
|--- Manufacturer_code <= 1.50
|   |--- mileage_mnorm <= 0.10
|   |   |--- Manufacturer_code <= 0.50
|   |   |   |--- class: 1
|   |   |--- Manufacturer_code > 0.50
|   |   |   |--- class: 0
|   |--- mileage_mnorm > 0.10
|   |   |--- Manufacturer_code <= 0.50
|   |   |   |--- class: 0
|   |   |--- Manufacturer_code > 0.50
|   |   |   |--- class: 0
|--- Manufacturer_code > 1.50
|   |--- Manufacturer_code <= 2.50
|   |   |--- mpg_mnorm <= 0.17
|   |   |   |--- class: 3
|   |   |--- mpg_mnorm > 0.17
|   |   |   |--- class: 3
|   |--- Manufacturer_code > 2.50
|   |   |--- Manufacturer_code <= 4.50
|   |   |   |--- class: 0
|   |   |--- Manufacturer_code > 4.50
|   |   |   |--- class: 3
```

```
# Decision Tree : Feature Importance
```

```
dtc_imp_features = pd.DataFrame({'feature': cars_inputs_names,
'importance': np.round(dtc_model.feature_importances_, 3)})
dtc_imp_features.sort_values('importance', ascending=False,
inplace=True); dtc_imp_features
```

	feature	importance
5	model_code	1.0
0	tax_mnorm	0.0
1	mileage_mnorm	0.0
2	mpg_mnorm	0.0
3	engineSize_mnorm	0.0
4	price_mnorm	0.0
6	transmission_code	0.0

```
7      fuelType_code      0.0
8  Manufacturer_code      0.0
```

```
# Decision Tree : Model Prediction (Training Subset)
```

```
dtc_model_predict = dtc_model.predict(cars_inputs_train);
dtc_model_predict
```

```
array([3, 3, 2, ..., 3, 0, 0], dtype=int32)
```

```
# Decision Tree : Prediction (Testing Subset)
```

```
dtc_predict = dtc_model.predict(cars_inputs_test); dtc_predict
```

```
array([3, 3, 1, ..., 1, 0, 0], dtype=int32)
```

```
# Decision Tree : Model Evaluation (Training Subset)
```

```
dtc_model_conf_mat = pd.DataFrame(confusion_matrix(cars_output_train,
dtc_model_predict)); dtc_model_conf_mat
dtc_model_perf = classification_report(cars_output_train,
dtc_model_predict); print(dtc_model_perf)
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	23628
1	1.00	1.00	1.00	16566
2	1.00	1.00	1.00	12079
3	1.00	1.00	1.00	25896
accuracy			1.00	78169
macro avg	1.00	1.00	1.00	78169
weighted avg	1.00	1.00	1.00	78169

```
import matplotlib.pyplot as plt
from sklearn.tree import plot_tree
```

```
# Set a larger figure size for better clarity
```

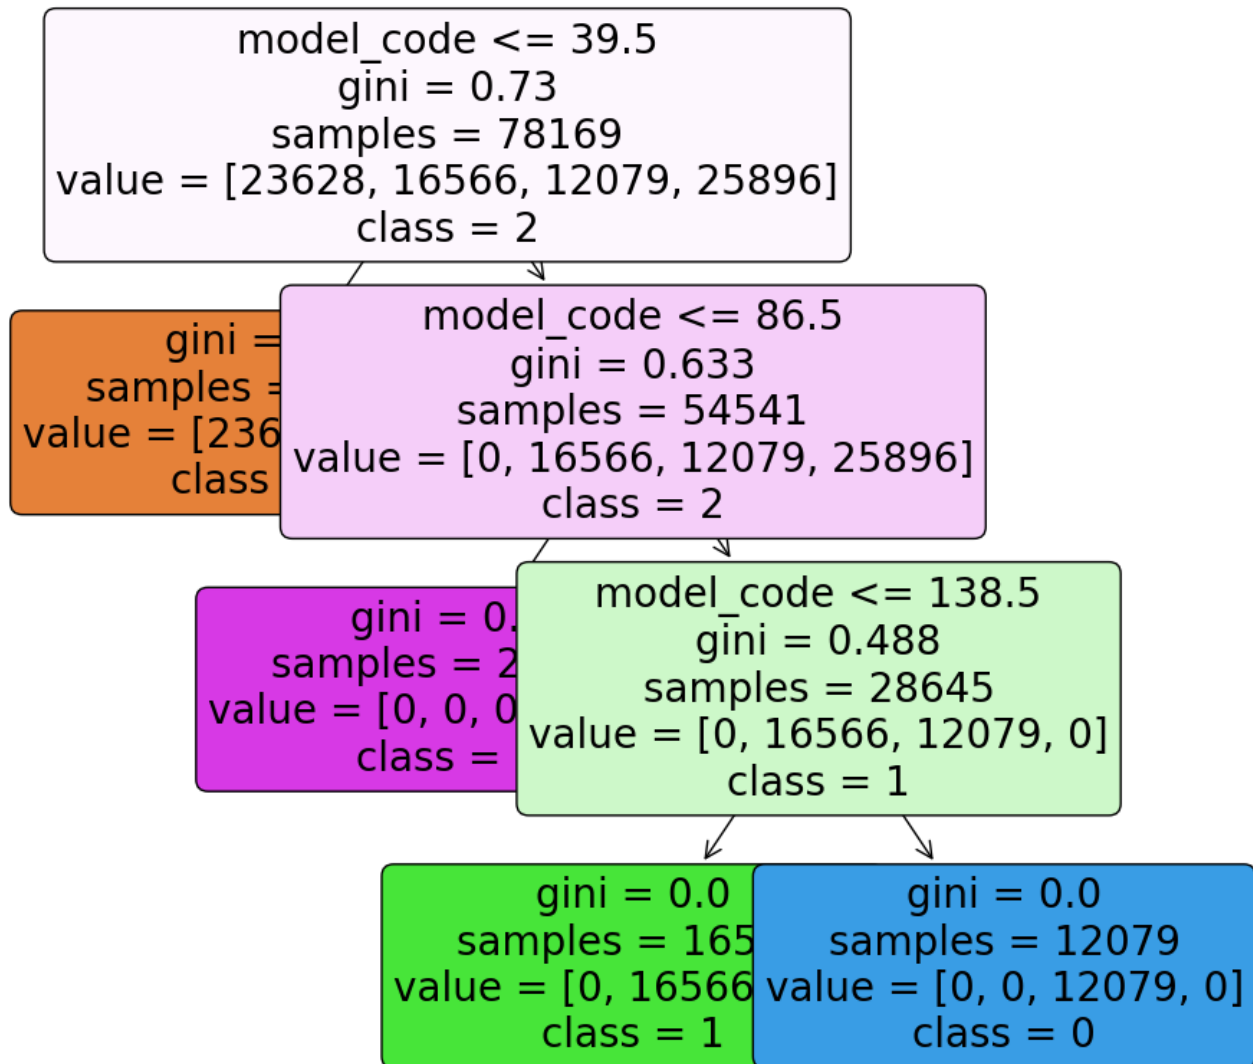
```
plt.figure(figsize=(10, 10))
```

```
# Plot the decision tree
```

```
train_subset_dtc_plot = plot_tree(dtc_model,
feature_names=cars_inputs_names, class_names=cars_output_labels,
rounded=True, filled=True, fontsize=20)
```

```
# Show the plot
```

```
plt.show()
```



```


# Set up the plot
ax = plt.axes()

# Plot the confusion matrix with annotations in integer format
sns.heatmap(dtc_model_conf_mat, annot=True, fmt='d', cmap='Paired')

# Set labels and title
ax.set_xlabel('Predicted Label')
ax.set_ylabel('True Label')
ax.set_title('Decision Tree : Confusion Matrix')

# Show the plot
plt.show()

```



```
# Cross Validation
from sklearn.model_selection import cross_val_score

# Define your decision tree classifier with desired parameters
dtc_cv = DecisionTreeClassifier(criterion='gini', random_state=45011)

# Perform 5-fold cross-validation
cv_scores = cross_val_score(dtc_cv, cars_inputs,
                             cars_output.values.ravel(), cv=20)
print("Cross-Validation Scores:", cv_scores)
print("Average Cross-Validation Score:", np.mean(cv_scores))

Cross-Validation Scores: [1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
Average Cross-Validation Score: 1.0

from sklearn.metrics import f1_score

# Compute F1 score
f1 = f1_score(cars_output_test, dtc_predict, average='macro') # or
'weighted' for weighted F1 score
print("F1 Score:", f1)

# Weighted F1 score
```



```

weighted_f1 = f1_score(cars_output_test, dtc_predict,
average='weighted')
print("Weighted F1 Score:", weighted_f1)

F1 Score: 1.0
Weighted F1 Score: 1.0

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split,
StratifiedShuffleSplit, cross_val_score
from sklearn.tree import DecisionTreeClassifier, export_text,
plot_tree
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix, f1_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import SelectFromModel
from sklearn.svm import SVC
import numpy as np
import time
import psutil

# Function to measure memory usage
def memory_usage():
    process = psutil.Process()
    return process.memory_info().rss / 1024 ** 2 # Memory usage in MB

# Start time
start_time = time.time()

# Data preprocessing and splitting
# Assuming you have your data loaded into cars_inputs and cars_output
cars_inputs_train, cars_inputs_test, cars_output_train,
cars_output_test = train_test_split(cars_inputs, cars_output,
test_size=0.2, random_state=42)

# End time
end_time = time.time()

# Time taken for data preprocessing and splitting
data_preprocessing_time = end_time - start_time

# Memory usage after data preprocessing
data_preprocessing_memory = memory_usage()

# Decision Tree
dt_start_time = time.time()
dt_model = DecisionTreeClassifier(criterion='gini',

```

```

random_state=45007, max_depth=3)
dt_model.fit(cars_inputs_train, cars_output_train)
dt_training_time = time.time() - dt_start_time
dt_memory_used = memory_usage()
dt_pred = dt_model.predict(cars_inputs_test)
dt_accuracy = accuracy_score(cars_output_test, dt_pred)

# Cross-validation for Decision Tree
dtc_cv_start_time = time.time()
dtc_cv = DecisionTreeClassifier(criterion='gini', random_state=45007)
cv_scores_dtc = cross_val_score(dtc_cv, cars_inputs,
cars_output.values.ravel(), cv=20)
dtc_cv_time = time.time() - dtc_cv_start_time
dtc_cv_accuracy = np.mean(cv_scores_dtc)

print("Decision Tree:")
print(f" - Training Time (s): {dt_training_time}")
print(f" - Memory Used (MB): {dt_memory_used}")
print(f" - Single Split Accuracy: {dt_accuracy}")
print(f" - Cross Validation Accuracy: {dtc_cv_accuracy}")
print()

Decision Tree:
  - Training Time (s): 0.05883908271789551
  - Memory Used (MB): 79.46875
  - Single Split Accuracy: 1.0
  - Cross Validation Accuracy: 1.0

```

## Random Forest

```

## Data Visualization Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.graph_objects as go
!pip install wordcloud
from wordcloud import WordCloud
from collections import Counter
from scipy import stats
from sklearn.tree import plot_tree
!pip install graphviz
import graphviz
from IPython.display import display
from collections import Counter

## Machine Learning Models and Evaluation Metrics

```

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.utils.validation import column_or_1d
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix, f1_score, precision_recall_fscore_support
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LogisticRegression, Lasso, Ridge
from sklearn.metrics import make_scorer
from sklearn.pipeline import make_pipeline
from sklearn.tree import export_graphviz
```

```
Requirement already satisfied: wordcloud in
./anaconda3/lib/python3.11/site-packages (1.9.3)
Requirement already satisfied: numpy>=1.6.1 in
./anaconda3/lib/python3.11/site-packages (from wordcloud) (1.24.3)
Requirement already satisfied: pillow in
./anaconda3/lib/python3.11/site-packages (from wordcloud) (9.4.0)
Requirement already satisfied: matplotlib in
./anaconda3/lib/python3.11/site-packages (from wordcloud) (3.7.1)
Requirement already satisfied: contourpy>=1.0.1 in
./anaconda3/lib/python3.11/site-packages (from matplotlib->wordcloud)
(1.0.5)
Requirement already satisfied: cycler>=0.10 in
./anaconda3/lib/python3.11/site-packages (from matplotlib->wordcloud)
(0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in
./anaconda3/lib/python3.11/site-packages (from matplotlib->wordcloud)
(4.25.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
./anaconda3/lib/python3.11/site-packages (from matplotlib->wordcloud)
(1.4.4)
Requirement already satisfied: packaging>=20.0 in
./anaconda3/lib/python3.11/site-packages (from matplotlib->wordcloud)
(23.0)
Requirement already satisfied: pyparsing>=2.3.1 in
./anaconda3/lib/python3.11/site-packages (from matplotlib->wordcloud)
(3.0.9)
Requirement already satisfied: python-dateutil>=2.7 in
./anaconda3/lib/python3.11/site-packages (from matplotlib->wordcloud)
(2.8.2)
Requirement already satisfied: six>=1.5 in
./anaconda3/lib/python3.11/site-packages (from python-dateutil>=2.7-
>matplotlib->wordcloud) (1.16.0)
```

```
Collecting graphviz
```

```
  Downloading graphviz-0.20.3-py3-none-any.whl (47 kB)
```

```
47.1/47.1 kB 2.8 MB/s eta
0:00:00
```

```
rf_classifier = RandomForestClassifier(n_estimators=100,
random_state=45050)
```

```

rf_classifier.fit(cars_inputs_train, cars_output_train['Cluster'])
RandomForestClassifier(random_state=45050)

y_train_pred_rf = rf_classifier.predict(cars_inputs_train)
y_test_pred_rf = rf_classifier.predict(cars_inputs_test)

# Train the Random Forest classifier
rf_classifier.fit(cars_inputs_train, cars_output_train['Cluster'])

# Print feature importances
feature_importances = rf_classifier.feature_importances_
feature_importance_df = pd.DataFrame({'Feature':
cars_inputs_train.columns, 'Importance': feature_importances})
sorted_feature_importance_df =
feature_importance_df.sort_values(by='Importance', ascending=False)
print("Feature Importances:")
print(sorted_feature_importance_df)

```

```

Feature Importances:

```

	Feature	Importance
5	model_code	0.796709
8	Manufacturer_code	0.107935
1	mileage_mmnorm	0.031623
2	mpg_mmnorm	0.030247
3	engineSize_mmnorm	0.015464
7	fuelType_code	0.005977
0	tax_mmnorm	0.005850
4	price_mmnorm	0.003353
6	transmission_code	0.002841

```

# For training set
print("Training Set Confusion Matrix:")
print(confusion_matrix(cars_output_train['Cluster'], y_train_pred_rf))

print("\nTraining Set Classification Report:")
print(classification_report(cars_output_train['Cluster'],
y_train_pred_rf))

```

```

Training Set Confusion Matrix:
[[23596  0  0  0]
 [  0 16567  0  0]
 [  0  0 12027  0]
 [  0  0  0 25979]]

```

```

Training Set Classification Report:

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	23596
1	1.00	1.00	1.00	16567
2	1.00	1.00	1.00	12027

3	1.00	1.00	1.00	25979
accuracy			1.00	78169
macro avg	1.00	1.00	1.00	78169
weighted avg	1.00	1.00	1.00	78169

```
# For testing set
print("\nTesting Set Confusion Matrix:")
print(confusion_matrix(cars_output_test['Cluster'], y_test_pred_rf))

print("\nTesting Set Classification Report:")
print(classification_report(cars_output_test['Cluster'],
y_test_pred_rf))
```

Testing Set Confusion Matrix:

```
[[5939  0  0  0]
 [  0 4141  0  0]
 [  0  0 3072  0]
 [  0  0  0 6391]]
```

Testing Set Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	5939
1	1.00	1.00	1.00	4141
2	1.00	1.00	1.00	3072
3	1.00	1.00	1.00	6391
accuracy			1.00	19543
macro avg	1.00	1.00	1.00	19543
weighted avg	1.00	1.00	1.00	19543

```
import pandas as pd
import numpy as np
import time
import psutil
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

# Function to measure memory usage
def memory_usage():
    process = psutil.Process()
    return process.memory_info().rss / 1024 ** 2 # Memory usage in MB

# Data preprocessing and splitting
# Assuming you have your data loaded into cars_inputs and cars_output
```

```

cars_inputs_train, cars_inputs_test, cars_output_train,
cars_output_test = train_test_split(cars_inputs, cars_output,
test_size=0.2, random_state=42)

# Initialize lists to store results
models = []
training_times = []
memory_used = []
single_split_accuracies = []
cross_validation_accuracies = []

# Decision Tree
dt_start_time = time.time()
dt_model = DecisionTreeClassifier(criterion='gini',
random_state=45050, max_depth=3)
dt_model.fit(cars_inputs_train, cars_output_train)
dt_training_time = time.time() - dt_start_time
dt_memory_used = memory_usage()
dt_pred = dt_model.predict(cars_inputs_test)
dt_accuracy = accuracy_score(cars_output_test, dt_pred)

# Cross-validation for Decision Tree
dtc_cv_start_time = time.time()
dtc_cv = DecisionTreeClassifier(criterion='gini', random_state=45050)
cv_scores_dtc = cross_val_score(dtc_cv, cars_inputs,
cars_output.values.ravel(), cv=20)
dtc_cv_time = time.time() - dtc_cv_start_time
dtc_cv_accuracy = np.mean(cv_scores_dtc)

# Append Decision Tree results to lists
models.append('Decision Tree')
training_times.append(dt_training_time)
memory_used.append(dt_memory_used)
single_split_accuracies.append(dt_accuracy)
cross_validation_accuracies.append(dtc_cv_accuracy)

# Random Forest
rf_start_time = time.time()
rf_classifier = RandomForestClassifier(n_estimators=100,
random_state=45011)
rf_classifier.fit(cars_inputs_train, cars_output_train['Cluster'])
rf_training_time = time.time() - rf_start_time
rf_memory_used = memory_usage()
y_train_pred_rf = rf_classifier.predict(cars_inputs_train)
y_test_pred_rf = rf_classifier.predict(cars_inputs_test)
rf_accuracy_train = accuracy_score(cars_output_train['Cluster'],
y_train_pred_rf)
rf_accuracy_test = accuracy_score(cars_output_test['Cluster'],
y_test_pred_rf)

```

```

# Append Random Forest results to lists
models.append('Random Forest')
training_times.append(rf_training_time)
memory_used.append(rf_memory_used)
single_split_accuracies.append(rf_accuracy_test) # Using test
accuracy as we already calculated it
cross_validation_accuracies.append(np.nan) # Cross-validation
accuracy not calculated here

# Create DataFrame
results_df = pd.DataFrame({
    'Model': models,
    'Training Time (s)': training_times,
    'Memory Used (MB)': memory_used,
    'Single Split Accuracy': single_split_accuracies,
    'Cross Validation Accuracy': cross_validation_accuracies
})

# Print DataFrame
print(results_df)

```

	Model	Training Time (s)	Memory Used (MB)	Single Split
Accuracy \				
0	Decision Tree	0.08861	86.921875	
1.0				
1	Random Forest	3.08000	75.968750	
1.0				
	Cross Validation Accuracy			
0		1.0		
1		NaN		